

GraphTerm: A notebook-like graphical terminal interface for collaboration and inline data visualization

Ramalingam Saravanan^{*†}



Abstract—The notebook interface, which blends text and graphics, has been in use for a number of years in commercial mathematical software and is now finding more widespread usage in scientific Python with the availability browser-based front-ends like the Sage and IPython notebooks. This paper describes a new open-source Python project, GraphTerm, that takes a slightly different approach to blending text and graphics to create a notebook-like interface. Rather than operating at the application level, it works at the unix shell level by extending the command line interface to incorporate elements of the graphical user interface. The XTerm terminal escape sequences are augmented to allow any program to interactively display inline graphics (or other HTML content) simply by writing to standard output.

GraphTerm is designed to be a drop-in replacement for the standard unix terminal, with additional features for multiplexing sessions and easy deployment in the cloud. The interface aims to be tablet-friendly, with features like clickable/tappable directory listings for navigating folders etc. The user can switch, as needed, between standard line-at-a-time shell mode and the notebook mode, where multiple lines of code are entered in cells, allowing for in-place editing and re-execution. Multiple users can share terminal sessions for collaborative computing.

GraphTerm is implemented in Python, using the Tornado web framework for the server component and HTML+Javascript for the browser client. This paper discusses the architecture and capabilities of GraphTerm, and provides usage examples such as inline data visualization using matplotlib and the notebook mode.

Index Terms—GUI, CLI, graphical user interface, command line interface, notebook interface, graphical shell

Introduction

Text and graphics form important components of the user interface when working with computers. Early personal computers only supported the textual user interface, more commonly known as the *command line interface* (CLI). However, when the Apple Macintosh popularized the *graphical user interface* (GUI), it soon became the preferred means for interacting with the computer. The GUI is more user-friendly, especially for beginners, and provides a more pleasant visual experience. The GUI typically provides buttons and widgets for the most common tasks, whereas the CLI requires recalling and typing out commands to accomplish tasks. However, the friendliness of the GUI comes at a cost—it can be much more difficult to perform advanced tasks using the GUI as compared to using

the CLI. Using a GUI is analogous to using a phrase book to express yourself in a foreign language, whereas using a CLI is like learning words to form new phrases in the foreign language. The former is more convenient for first-time and casual users, whereas the latter provides the versatility required by more advanced users.

The dichotomy between the textual and graphical modes of interaction also extends to scientific data analysis tools. Traditionally, commands for data analysis were typed into a terminal window with an interactive shell and the graphical output was displayed in a separate window. Some commercial software, such as Mathematica and Maple, provided a more integrated notebook interface that blended text and graphics, thus combining aspects of the CLI with the GUI. One of the exciting recent developments in scientific Python has been the development of alternative, open source, notebook interfaces for scientific computing and data analysis—the Sage and IPython notebooks [Perez12]. Since Python is a more general-purpose language than Mathematica or Maple, the notebook interface could potentially reach a much wider audience.

A notebook display consists of a sequence of cells, each of which can contain code, figures, or text (with markup). Although originally developed for exploratory research, notebooks can be very useful for presentations and teaching as well. They can provide step-by-step documentation of complex tasks and can easily be shared. The cells in a notebook do not necessarily have to be executed in the sequence in which they appear. In this respect, the notebook interface can be considered an expression of “literate programming”, where snippets of code are embedded in natural language documentation that explains what the code does [Knuth84].

Another emerging area where the notebook interface could serve as an important tool is *reproducible research* [Stodden13]. As computational techniques are increasingly being used in all areas of research, reproducing a research finding requires not just the broad outline of the research methodology but also documentation of the software development environment used for the study. The need for reproducible research is highlighted by the recent controversy surrounding the highly influential Reinhart-Rogoff study that identified a negative relationship between a country’s debt and its economic growth rate. A follow-up study [Herndon13] identified a simple coding error that affects key findings of the original study. The self-documenting nature of code and results presented in a notebook format can make it easy to

^{*} Corresponding author: sarava@tamu.edu

[†] Texas A&M University

share and reproduce such computations.

Background

The author had some experience with commercial notebook interfaces before, but used the IPython Notebook interface for the first time in January 2013, when teaching an introductory undergraduate programming course for geoscientists using Python. After initially using the command line Python interpreter, the class switched to using IPython Notebook, whose inline code editing and graphics display turned out to be really convenient. The notebook interface was used for presenting lecture material, and the students used it for their programming assignments, turning in their notebooks for grading (in PDF format).

The author had previously been working on a project called GraphTerm, which implements a “graphical terminal interface” using a Python backend and a HTML5+Javascript frontend [GraphTerm]. It was a follow-up to two earlier projects, the browser-based AjaxTerm, and XMLTerm, a GUI-like browser built using the Mozilla framework [Sarava00]. GraphTerm is aimed at being a drop-in replacement for XTerm, the standard unix terminal, with additional graphical and collaborative features. It retains all the features of the CLI, including pipes, wildcards, command recall, tab completion etc., and also incorporates web-based sharing, as well as GUI-like features, such as clickable folder navigation, draggable files, inline image display etc. (There also other terminal projects with similar goals, such as TermKit for OS X and Terminology for Linux.)

The distinctive features of the notebook interface, such as inline editing and graphics, are not specific to any particular programming language or interactive shell. Also, the GraphTerm code already had the capability to incorporate GUI-like features into the terminal. Therefore, it seemed worth experimenting with GraphTerm to see how far it could be extended to support a generic, language-independent, notebook interface, while still retaining *full backward compatibility* with the unix terminal. The goal was to allow the terminal to be switched to a notebook mode, regardless of what application was running in the shell. The backward compatibility requirements and the loose coupling between the notebook and the underlying application could make it more fragile and restricted, but that would be an unavoidable trade-off. The rest of this paper reports on the results of this effort to combine the CLI, GUI, and the notebook interface.

Implementation

The standard unix terminal supports two types of buffers: (i) the normal *scroll buffer* that contains lines of text, and (ii) the *full screen buffer* used by text editors like `vi` etc. Special character strings known as *escape sequences* are output by programs to switch the terminal between the two buffers [XTerm]. GraphTerm currently supports most of the standard XTerm escape sequences and introduces additional escape sequences that allow display of HTML fragments in the scroll buffer and the full screen buffer. The HTML fragments can




Fig. 1: Architecture of GraphTerm. Browser client connects to Tornado server using websockets. Hosts connect to server using TCP.

contain just about anything that can be displayed on a web page, including text with markup, tables, and images.

The GraphTerm server is written in pure python, using the [Tornado web framework](#), with websocket support. The browser client uses standard HTML5+Javascript+CSS (with jQuery). The code is released under the BSD License and the repository is available on [Github](#).

The GraphTerm server may be run on the desktop or on a remote computer. Users create and access terminal sessions by the connecting to the Graphterm server on the default port 8900, either directly or through SSH port forwarding (Figure 1). By default, the localhost on the computer where the GraphTerm server is running is available for opening terminal sessions. Other computers can also connect to the GraphTerm server, on a different port (8899), to make them accessible as hosts for connection from the browser.

A pseudo-tty (`pty`) device is opened on the host for each terminal session. By setting the `PROMPT_COMMAND` environment variable, GraphTerm determines when the standard output of the previous command ends, and the prompt for the new command begins. The connection between the browser and the GraphTerm server is implemented using websockets (bi-directional HTTP). The GraphTerm server acts as a router sending input from controlling browser terminal sessions to the appropriate `pty` on the host computer, and transmitting output from each `pty` to all connected browser terminal sessions.

All the scroll buffer and full screen buffer content is stored on the server, which means that the terminal is persistent across different browser sessions. For example, you can leave the terminal on your desktop computer at work and access the exact same content on your laptop browser when you get home. This allows GraphTerm to be used like the GNU `screen` or `tmux` programs. Storing the content on the server also allows multiple users to share access to the same terminal session for collaboration, similar to, e.g., Google Docs. This means that multiple users will be able to view and modify a GraphTerm notebook session in real time.

```
graphterm$ sh helloworld.sh
Hello World!
```



```
graphterm$
```

Fig. 2: Output of `helloworld.sh` within *GraphTerm*, showing inline HTML text and image.

The GraphTerm API

Programs running within a *GraphTerm* shell communicate with it by writing to its standard output a block of text using a format similar to a HTTP response, preceded and followed by XTerm-like escape sequences:

```
\x1b[?1155;<cookie>h
{"content_type": "text/html", ...}

<div>
...
</div>
\x1b[?1155l
```

where `<cookie>` denotes a numeric value stored in the environment variable `GTERM_COOKIE`. This random cookie is a security measure that prevents malicious files from accessing *GraphTerm*. The opening escape sequence is followed by an *optional* dictionary of header names and values, using JSON format. This is followed by a blank line, and then any data (such as the HTML fragment to be displayed).

A simple bash shell script, `hello_world.sh`, illustrates this API:

```
#!/bin/bash
# A Hello World program using the GraphTerm API

prefix=https://raw.githubusercontent.com/mitotic/graphterm
url=$prefix/master/graphterm/www/GTTY500.png
esc='printf "\033\'
code="1155"
# Prefix escape sequence
echo "${esc}[?${code}h"
# Display text with HTML markup
echo '<b>Hello</b>'
echo '<b style="color: red;">World!</b><p>'
# Display inline image
echo "<a><img width='200' src='\"$url\"'></a>"
# Suffix escape sequence
echo "${esc}[?${code}l"
```

If run within *GraphTerm*, the script produces the output shown in Figure 2.

Features

GraphTerm is written in pure Python and the only dependency is the `tornado` web server module. It can be installed using `easy_install` or `setuptools`. Once the *GraphTerm* server program is started, it listens on port 8900 on `localhost` by default, and any browser can be used to connect to it and open new terminal sessions using the URL `http://localhost:8900`. At this point, *GraphTerm* can be used like a regular terminal, with commands like `ls`,

```
scipy_proceedings$ cd papers/
papers$ ls
00_vanderwalt  r_saravanan
papers$ gls
..             .             ~
00_vanderwalt  r_saravanan
papers$
```

Fig. 3: Output of `ls` and `gls` commands for the same directory. The names displayed by `gls` are hyperlinked, and may be clicked to navigate to a folder or open a file.

```
papers$ gmenu icons
papers$ gls
```

```
papers$
papers$ cd 00_vanderwalt; gls -f
```

```
00_vanderwalt.rst  fig1.png
00_vanderwalt$
```

Fig. 4: Output of `gls` with icon display enabled. Clicking on the folder icon for `00_vanderwalt` (red rectangle) executes the command `cd 00_vanderwalt; gls -f` via the command line (green rectangle) to navigate to the folder and list its directory contents. (This action also overwrites any immediate previous file navigation command in the *GraphTerm* command history, to avoid command clutter.)

`vi`, etc. However, to use the graphical capabilities of *GraphTerm*, one needs to use *GraphTerm*-aware versions of these commands, with names like `gls` and `gvi`, that are part of the command toolchain that is bundled with the code. The toolchain commands communicate using pipes and may be written any language, e.g., Bash shell script, Python etc., using the API described above. The GUI-like features of *GraphTerm* implemented using this toolchain are discussed and illustrated below.

Clickable folders and files

The output of the standard `ls` command displays the directory listing as plain text, whereas the `gls` command from the toolchain displays a hyperlinked (“clickable”) directory listing (Figure 3).

By default, `gls` does not display icons or images in the directory listing. However, icon display can be enabled using the *GraphTerm* menubar (Figure 4).

You can navigate folders in *GraphTerm* using GUI-like actions, like you would do in the Windows Explorer or the

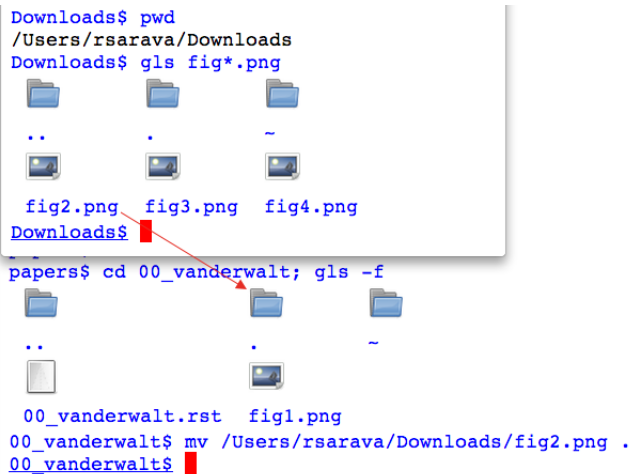


Fig. 5: File `fig2.png` is dragged from the Downloads folder from the source terminal and dropped into the `.` (current directory) folder icon displayed by `gls` in the destination terminal. This executes the command `mv /user/rsarava/Downloads/fig2.png .` in the destination terminal to move the file.

Mac Finder, while retaining the ability to drop back to the CLI at any time. If the current command line is empty, clicking on a hyperlinked folder will insert a new command line of the form:

```
cd newdir; gls -f
```

which will change the current directory to `newdir` and list its contents. Clicking on a hyperlinked filename will generate a new command line to invoke platform-dependent commands like `open` or `xdg-open` to open the file using the default program for its file type. This feature illustrates one of the basic design goals of GraphTerm, that each GUI-like action should generate a corresponding shell command that actually carries out that action. This allows the action to be logged and reproduced later.

Drag and drop

GraphTerm currently provides limited support for drag-and-drop operations, including support for uploading/copying files between terminal sessions on different computers connected to the same GraphTerm server. As shown in Figure 5, when a file is dragged from the source terminal and dropped into a folder displayed in the destination terminal, a `mv` command is generated to perform the task. Thus the GUI action is recorded in the command line for future reference.

Session sharing and theming

GraphTerm terminal sessions can be shared between multiple computers, with different types of access levels for additional users accessing the same terminal, such as read-only access or full read-write access. Since a GraphTerm terminal session is just a web page, it also supports theming using CSS stylesheets. The terminal sharing and theming are decoupled, which means that two users can view the same terminal using different themes (Figure 6)!

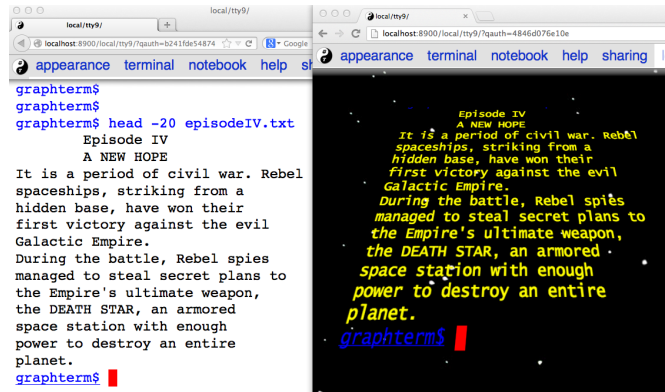


Fig. 6: Two shared views of a GraphTerm terminal session showing the output of the command `head -20 episodeIV.txt` on a computer running OS X Lion. The left view is in a Firefox window with the **default** theme and the right view shows the same terminal in a Chrome window, using the **stars3D** perspective theme (which currently does not work on Firefox).

Inline graphics

Since GraphTerm can display arbitrary HTML fragments, it is easy to display graphical output from programs. The `gimage` command in the toolchain can be used to display inline images. The toolchain also includes the `yweather` command to display the current weather forecast graphically using the Yahoo Weather API. Other toolchain commands include `glandslide` to use the Python-based `landslide` presentation tool and `greveal` that uses `reveal.js` to display slideshows within a GraphTerm window.

GraphTerm can be used for inline display of graphical output from `matplotlib` (Figure 7). The API bundled with GraphTerm uses the `StringIO` module to capture the binary plot data using the `png` image output produced by the `Agg` renderer and then displays the image using GraphTerm escape sequences. A module called `gmatplot` is supplied with GraphTerm to provide explicit access to this plotting API. Another module `gpylab` is also provided, for *monkey patching* existing plotting code to work within GraphTerm with little or no changes. For example, if the Python interpreter is invoked using the following command:

```
python -i $GTERM_DIR/bin/gpylab.py
```

then `pylab` functions like `draw`, `figure`, and `show` will automatically use the GraphTerm API to display inline graphics (e.g. see the notebook example shown in Figure 8).

Since communication with GraphTerm occurs solely via the standard output of a program, inline graphics can be displayed from any plotting program, including commercial software like IDL and other plotting packages like the NCAR Command Language (NCL). Inline graphics display can also be used across SSH login boundaries by including support for the GraphTerm API in the plotting program on the remote machine.

Notebook mode

GraphTerm can be switched from the normal terminal mode to a blank notebook mode using the key sequence `Shift-Enter`


```
bin$ python exercise-viz-contour-gterm.py
```




Fig. 7: Inline display of a 2-dimensional filled contour plot of surface air temperature on the globe, generated by `matplotlib`. The code for this plot is taken from the textbook by [Lin12].

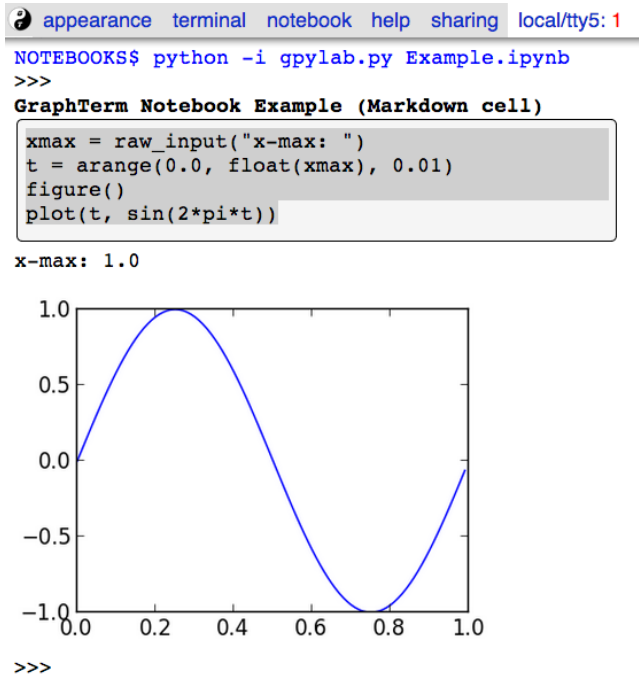


Fig. 8: GraphTerm notebook mode, where the notebook contents are read from a file saved using the `ipynb` format. The first cell contains Markdown text and the second cell contains python code to generate a simple plot using `matplotlib`. Note the use of `raw_input` to prompt the user for terminal input.

or using the menubar. The user can also click on a notebook file displayed in the `gls` directory listing to open it and pre-fill the notebook cells with content from the file (Figure 8). The notebook mode supports the normal terminal operations, such as reading from the standard input (i.e., `raw_input` in Python) and using debuggers, as well as GraphTerm extensions like inline graphics. (Full screen terminal operations are not supported in the notebook mode.)

Users can save the contents of the displayed notebook to a file at any time. Users exit the notebook mode and revert

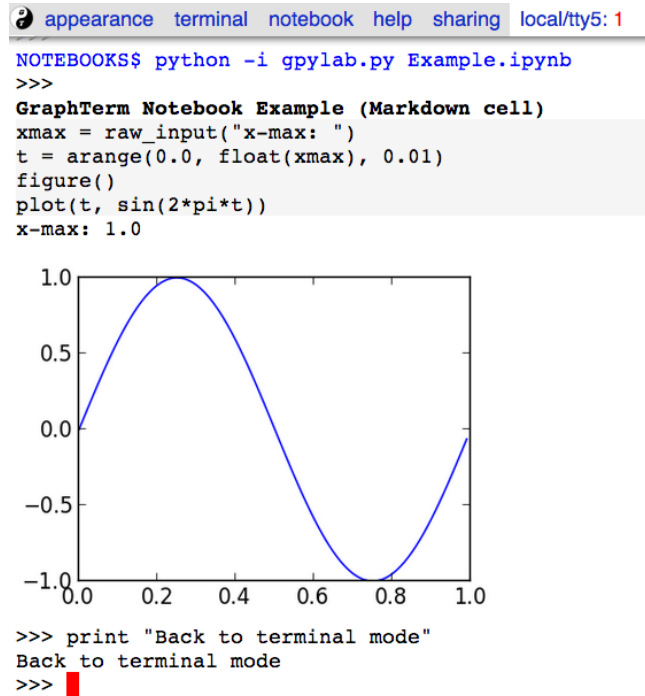


Fig. 9: When switching back to the terminal mode after exiting the notebook mode, the notebook contents can either be discarded or be appended like normal terminal output, as shown above.

to the normal terminal mode using the menubar or simply by typing *Control-C*. When exiting the notebook mode, users can choose to either merge all the notebook content back into the terminal session or discard it (Figure 9).

The notebook implementation in GraphTerm attempts to preserve interoperability with the IPython Notebook to the extent possible. GraphTerm can read and write notebooks using the IPython Notebook format (`*.ipynb`), although it uses the [Markdown](#) format for saving notebook content. (Markdown was chosen as the native format because it is more human-friendly than ReStructuredText or JSON, allows easy concatenation or splitting of notebook files, and can be processed by numerous Markdown-aware publishing and presentation programs like `landslide` and `reveal.js`.) GraphTerm supports many of the same keyboard shortcuts as IPython Notebook. GraphTerm can also be used with the command-line version of IPython. However, the generic, loosely-coupled notebook interface supported by GraphTerm will never be able to support all the features of IPython Notebook.

Here is how the notebook mode is implemented within GraphTerm: when the user switches to the notebook mode, a separate scroll buffer is created for each cell. When the user executes a line of code within a GraphTerm notebook cell, the code output is parsed for prompts to decide whether to continue to display the output in the output cell, or to return focus to the input cell. This text-parsing approach does make the GraphTerm notebook implementation somewhat fragile, compared to other notebook implementations that have a tighter coupling with the underlying code interpreter (or kernel). However it allows GraphTerm to work with interactive

view terminal notebook help share local/tty1:1 steal run R-ggplot.R

notebooks\$ R -q # Notebook: R-ggplot.R.md

>

GraphTerm Notebook using R: Example 2

```
# Load GraphTerm API helper functions
source(paste(Sys.getenv("GTERM_DIR"), "/bin/gtermapi.R",
  sep=""))
```

>

```
g <- gcairo() # Setup device for GraphTerm
require("ggplot2")
p <- ggplot(mtcars, aes(factor(cyl), mpg))
p + geom_boxplot()
g$frame() # Display plot as inline image
```

Loading required package: ggplot2
format = ARGB (400 x 300)




Fig. 10: Inline graphics in notebook mode when running the standard R interpreter within GraphTerm.

shells for any platform, such as R (Figure 10) (or any interactive program with prompts, including closed source binaries for languages like IDL).

Since all GraphTerm content is stored on the server, the notebook can be accessed by multiple users simultaneously for collaboration. Like inline graphics, the notebook mode works transparently when executing interactive shells after a remote SSH login, because all communication takes place via the standard output of the shell. The non-graphical notebook mode can be used without the remote program ever being aware of the notebook interface. However, the remote program will need to use the GraphTerm escape sequences to display inline graphics within the notebook.

Conclusion

The GraphTerm project extends the standard unix terminal to support many GUI-like capabilities, including inline graphics display for data analysis and visualization. Adding features like clickable folder navigation to the CLI also makes it more touch-friendly, which is likely to be very useful on tablet computers. Incorporating GUI actions within the CLI allows recording of many user actions as scriptable commands, facilitating reproducibility.

GraphTerm also demonstrates that the notebook interface can be implemented as an extension of the CLI, by parsing the textual output from interactive shells. This allows the notebook interface to be “bolted on” to any interactive shell program and to be used seamlessly even across SSH login boundaries. The notebook features and the real-time session sharing capabilities could make GraphTerm an useful tool for collaborative computing and research.

REFERENCES

- [GraphTerm] *GraphTerm home page* <http://code.mindmeldr.com/graphterm>
- [Herndon13] T. Herndon, M. Ash, and R. Pollin. *Does High Public Debt Consistently Stifle Economic Growth? A Critique of Reinhart and Rogoff* http://www.peri.umass.edu/fileadmin/pdf/working_papers/working_papers_301-350/WP322.pdf
- [Knuth84] D. Knuth. *Literate Programming*. The Computer Journal archive. Vol. 27 No. 2, May 1984, pp. 97-111 <http://literateprogramming.com/knuthweb.pdf>
- [Lin12] J. Lin. *A Hands-On Introduction to Using Python in the Atmospheric and Oceanic Sciences* [Chapter 9, Exercise 29, p. 162] <http://www.johnny-lin.com/pyintro>
- [Perez12] F. Perez. *The IPython notebook: a historical retrospective*. Jan 2012 <http://blog.fperez.org/2012/01/ipython-notebook-historical.html>
- [Sarava00] R. Saravanan. *XMLterm: A Mozilla-based Semantic User Interface*. XML.com, June 2000 <http://www.xml.com/pub/a/2000/06/07/xmlterm/>
- [Stodden13] V. Stodden, D. H. Bailey, J. Borwein, R. J. LeVeque, W. Rider, and W. Stein. *Setting the Default to Reproducible: Reproducibility in Computational and Experimental Mathematics*. February 2013 http://stodden.net/icerm_report.pdf
- [XTerm] *XTerm Control Sequences* <http://invisible-island.net/xterm/ctlseqs/ctlseqs.html>